

Tran Quoc Truong

Applied AI Engineer · AI Agents & RAG · Backend Systems

📍 Ho Chi Minh City, Vietnam

✉ tranquoctruong20@gmail.com

🌐 portfolio.truong51972.id.vn

🌐 truong51972

🌐 truong51972

Summary

Applied AI Engineer focused on translating business and operational challenges into practical, reliable AI solutions. Combines solution architecture and hands-on engineering to address enterprise knowledge access, intelligent automation, and AI agent reliability, with end-to-end ownership from problem framing and architectural decisions to delivery and quality validation.

Technical Skills

AI & Agentic Systems: RAG, Hybrid Retrieval, LangGraph, MCP (Model Context Protocol), Tool Calling, LiteLLM, Qdrant, Docling

ML, Fine-Tuning & Model Serving: PyTorch, Hugging Face Transformers, PEFT, Unsloth, LoRA/QLoRA, vLLM, Synthetic Data Generation

LLM Evaluation & Reliability: OpenTelemetry, Runtime Trace Reconstruction, LLM-as-a-Judge, Quality Gates, Guardrails, Langfuse

Backend Engineering: Python, FastAPI, Django/DRF, Pydantic, SQLAlchemy, PostgreSQL, Redis, Celery, SSE / Streaming APIs

Test Automation & Quality Engineering: Playwright, Cypress, Katalon Studio, Postman, API Testing, Regression Testing, SQL Validation

Infrastructure & Delivery: Docker, Docker Compose, RunPod, Nginx, MinIO, CI/CD, Linux

Experience

FPT Software , Applied AI Engineer & Automation Tester (IVS)	Ho Chi Minh City, Vietnam
IQP (Code Intelligence): Designed multi-repository code intelligence combining semantic, lexical, and code-graph retrieval to provide traceable implementation and test context through MCP-enabled AI workflows. Agent Assurance: Architected trace-based evaluation that reconstructs agent execution evidence against versioned requirements, supporting auditable regression analysis and governed quality-gate decisions. Banking Automation: Reduced manual regression effort by ~70% across banking and payment workflows using reusable API, database, and cross-platform automation with SQL-based transaction validation.	Dec 2024 – present 1 year 11 months
FPT Software , Full-Stack Developer Intern (FSA)	Ho Chi Minh City, Vietnam
- Built the core backend for an online learning and assessment platform using Django with role-based access control and automated exam workflows. - Integrated a separate FastAPI computer-vision service for AI-assisted real-time exam proctoring and face detection. - Containerized and deployed the multi-service architecture using Docker Compose, Nginx reverse proxy, and Cloudflare Tunnel.	Sept 2024 – Dec 2024 4 months

Achievements

Golden Solution Prize (1st Place) — IVS Solution Day 2.0 2026	Sept 2026
Recognized for Flezi Polaris, an Agent Assurance solution combining execution traceability, behavioral evaluation, and auditable quality gates.	
2nd Runner-up (3rd Place) — IVS Hackathon 2026	Aug 2026
Independently architected the Agent Assurance engine later integrated into Omni-Agent, enabling trace-based evaluation and quality verification.	
Second Prize — Autonomous Driving Research Paper Competition	Apr 2024
Sole-authored the project, from dataset labeling and YOLOv8 training to lane segmentation and PID steering control.	

Featured Projects

Omni-Agent — Applied AI Platform (Knowledge Builder & Agent Assurance)

Independently architected and built a pre-production Applied AI platform for enterprise document intelligence, grounded multi-agent workflows, and reproducible AI-agent assurance.

- Solution Architecture: Designed project-scoped FastAPI services with independent data ownership, authorization boundaries, and durable asynchronous workflows.
- Knowledge Engineering: Built versioned Docling/OCR processing and PostgreSQL–Qdrant hybrid retrieval, preserving usable documents through recoverable indexing failures.
- Agentic RAG & Trust: Orchestrated LangGraph research agents through MCP, with run-scoped citations, immutable answer evidence, and input/context/output guardrails.
- Agent Assurance: Built versioned Scenarios and qualified Run snapshots, linking OpenTelemetry evidence with deterministic quality gates for reproducible agent evaluation.
- Product Engineering: Reduced first-test STFV from 33 to 6 actions (4 for demo) through guided onboarding in a controlled Playwright fixture, preserving preflight checks and informed consent.

APIT — Agent Programmatic Integration Testing

AI-assisted API testing framework generating structured test scenarios from API documentation via fine-tuned LLMs.

- Developed a bilingual synthetic dataset generation pipeline producing 1,258 domain-specific API testing samples.
- Fine-tuned Qwen2.5-3B with LoRA/QLoRA, achieving macro-F1 of 0.655 (+41.5% relative improvement over Llama-3.2-3B baseline).
- Deployed the fine-tuned model on RunPod through an OpenAI-compatible vLLM service with quantized inference and runtime LoRA adapters.
- Implemented a hybrid evaluation pipeline combining LLM-as-a-Judge, fuzzy matching, and deterministic JSON-schema validation.

E-Commerce AI Assistant — Conversational RAG & Search

End-to-end shopping assistant combining filtered vector search and multi-agent conversational retrieval.

- Built filtered semantic product retrieval with Milvus and conversational consultation through LangGraph/LangChain.
- Delivered a Docker-based prototype integrating FastAPI, PostgreSQL, Redis, and vector search.

Education

FPT School of Business & Technology , Master of Software Engineering in AI	Ho Chi Minh City, Vietnam May 2026 – present
FPT University , Bachelor of Artificial Intelligence	Ho Chi Minh City, Vietnam 2021 – 2025

Languages

- Vietnamese:** Native
- English:** Professional working proficiency